

Session 4: Tidying data Confidence Intervals

Kostis Christodoulou
London Business School



Tidy data (1/3)

The tables below show the same values of four variables **country**, **year**, **cases**, **population**, but each dataset organises the values in a different way.

```
> table1
# A tibble: 6 x 4
  country    year cases population
  <chr>    <int> <int>    <int>
1 Afghanistan 1999     745 19987071
2 Afghanistan 2000    2666 20595360
3 Brazil      1999   37737 172006362
4 Brazil      2000   80488 174504898
5 China       1999  212258 1272915272
6 China       2000  213766 1280428583
> table2
# A tibble: 12 x 4
  country    year type      count
  <chr>    <int> <chr>    <int>
1 Afghanistan 1999 cases      745
2 Afghanistan 1999 population 19987071
3 Afghanistan 2000 cases      2666
4 Afghanistan 2000 population 20595360
5 Brazil      1999 cases   37737
6 Brazil      1999 population 172006362
7 Brazil      2000 cases   80488
8 Brazil      2000 population 174504898
9 China       1999 cases   212258
10 China      1999 population 1272915272
11 China      2000 cases   213766
12 China      2000 population 1280428583
```


```
> table3
# A tibble: 6 x 3
  country    year rate
  <chr>    <int> <chr>
1 Afghanistan 1999 745/19987071
2 Afghanistan 2000 2666/20595360
3 Brazil      1999 37737/172006362
4 Brazil      2000 80488/174504898
5 China       1999 212258/1272915272
6 China       2000 213766/1280428583
> table4a
# A tibble: 3 x 3
  country    `1999` `2000`
  <chr>    <int> <int>
1 Afghanistan     745     2666
2 Brazil        37737    80488
3 China         212258   213766
> table4b
# A tibble: 3 x 3
  country    `1999` `2000`
  <chr>    <int> <int>
1 Afghanistan 19987071 20595360
2 Brazil      172006362 174504898
3 China       1272915272 1280428583
```

Tidy data 2/3

Tidy data is a specific way of organising data to facilitate analysis with the tidyverse. The tidy data standard has been designed to facilitate exploratory data analysis; tidy datasets and tidy tools help make data analysis easier, allowing you to focus on the interesting domain problem, not on the logistics of cleaning data.

There are three rules which make a dataset tidy:

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.



We often need to reshape our datasets and should have a way to go:

- from wide format to long (tidy) format using *gather()* or *pivot_longer()*
- from long (tidy) to wide format using *spread()* or *pivot_wider()*

Tidy data 3/3

Five main ways data tend not to be tidy:

1. Column headers are values, not variable names.
2. Multiple variables are stored in one column.
3. Variables are stored in both rows and columns.
4. Multiple types of observational units are stored in the same table.
5. A single observational unit is stored in multiple tables.

```
> table1
# A tibble: 6 x 4
  country    year cases population
  <chr>    <int> <int>    <int>
1 Afghanistan 1999     745 19987071
2 Afghanistan 2000    2666 20595360
3 Brazil      1999   37737 172006362
4 Brazil      2000   80488 174504898
5 China       1999  212258 1272915272
6 China       2000  213766 1280428583

> table2
# A tibble: 12 x 4
  country    year type      count
  <chr>    <int> <chr>    <int>
1 Afghanistan 1999 cases        745
2 Afghanistan 1999 population 19987071
3 Afghanistan 2000 cases        2666
4 Afghanistan 2000 population 20595360
5 Brazil      1999 cases       37737
6 Brazil      1999 population 172006362
7 Brazil      2000 cases       80488
8 Brazil      2000 population 174504898
9 China       1999 cases      212258
10 China      1999 population 1272915272
11 China      2000 cases      213766
12 China      2000 population 1280428583
```

```
> table3
# A tibble: 6 x 3
  country    year rate
  <chr>    <int> <chr>
1 Afghanistan 1999 745/19987071
2 Afghanistan 2000 2666/20595360
3 Brazil      1999 37737/172006362
4 Brazil      2000 80488/174504898
5 China       1999 212258/1272915272
6 China       2000 213766/1280428583

> table4a
# A tibble: 3 x 3
  country `1999` `2000`
  <chr>    <int> <int>
1 Afghanistan 745    2666
2 Brazil      37737 80488
3 China      212258 213766

> table4b
# A tibble: 3 x 3
  country    `1999`    `2000`
  <chr>    <int>    <int>
1 Afghanistan 19987071 20595360
2 Brazil      172006362 174504898
3 China      1272915272 1280428583
```

Why tidy? (and not dirty/untidy/messy)

1. There's a general advantage to picking one consistent way of storing data.
2. Having variables in columns allows to exploit R's vectorised nature, as most built-in R functions work with vectors of values. That makes transforming tidy data feel particularly natural.

```
# The easiest way to get tidyr is to install the whole tidyverse:
install.packages("tidyverse")

# Alternatively, install just tidyr:
install.packages("tidyr")

# Or the development version from GitHub:
# install.packages("devtools")
devtools::install_github("tidyverse/tidyr")
```

Use the development version of ***tidyr*** to use
pivot_longer() and ***pivot_wider()***

pivot_longer() or gather()

- A common problem is when column names are not **names** of variables, but **values** of a variable.
- Column names **1999** and **2000** represent values of the year variable, and each row represents two observations, not one.

```
> table4a
# A tibble: 3 x 3
  country    `1999`    `2000`
*   <chr>      <int>    <int>
1 Afghanistan    745      2666
2 Brazil        37737    80488
3 China         212258    213766
```

We need to gather those columns into a new pair of variables. We need to map the names and the values.

1. The set of columns that represent values, not variables. In this example, those are the columns **1999** and **2000**.
2. The name, or key, of the variable whose values will form the column names: **year**.
3. The name of the variable whose values are spread over the cells: **cases**

```
table4a %>%
  pivot_longer(cols=c(`1999`, `2000`), names_to = "year", values_to = "cases")

table4a %>%
  gather(`1999`, `2000`, key = "year", value = "cases")

# A tibble: 6 x 3
  country    year    cases
  <chr>    <chr>    <int>
1 Afghanistan 1999      745
2 Brazil      1999    37737
3 China       1999   212258
4 Afghanistan 2000     2666
5 Brazil      2000    80488
6 china       2000   213766
```

Wide and long data (1/3)

There are three rules which make a dataframe tidy:

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.

Is this data in tidy format? What are the variables and observations?

student_id	final_exam	midterm	group_project
2457625	79	68	71
1758293	92	73	67
1622247	71	87	74

Work with your neighbour and sketch out on a piece of paper what this data would look like in tidy format. Some hints:

1. What are the variables?
2. What variable pairings are of interest?
3. Is each observation in its own row ?

Wide and long data (3/3)

We often need to reshape our datasets and we have a way to go:

- from wide format to long (tidy) format using *pivot_longer()* or *gather()*
- from long (tidy) to wide format using *pivot_wider()* or *spread()*

```
pivot_longer(cols=c('final_exam', 'midterm', 'group_project'), names_to = "assignment", values_to = "score")
gather('final_exam', 'midterm', 'group_project', key = assignment, value = score)
wide
```

id	x	y	z
1	a	c	e
2	b	d	f

Tidying data

id	x	y	z
1	a	c	e
2	b	d	f

```

pivot_longer(cols=c('final_exam', 'midterm', 'group_project'), names_to = "assignment", values_to = "score")
gather('final_exam', 'midterm', 'group_project', key = assignment, value = score)
pivot_wider(names_from = assignment, values_from = score)
spread(key = assignment, value = score)

```

Worked Examples: Live Coding